

Attack against 6th Byte

December 6, 2022

1 Complete Attack to Recover the 6th Byte of the Key

This notebook shows a complete attack to a single byte, specifically the 6th byte (index 5).

The train-configuration is composed by 2 devices (D1, D2) and all possible keys (all device-key permutations), except K0, which is attacked.

The attack-configuration is D3-K0.

The scenario is Multiple Device Model (MDM), where multiple devices are used in training and a different device is attacked.

The following results and metrics are shown: * Train Loss & Validation Loss * Train Accuracy & Validation Accuracy * Attack Loss & Attack Accuracy * Confusion Matrix * Guessing Entropy

```
[1]: # Basics
import json
import numpy as np
import polars as pl
import matplotlib.pyplot as plt
from tqdm import tqdm
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model
from sklearn.preprocessing import StandardScaler

# Custom
import sys
sys.path.insert(0, '../src/utils')
import helpers
import results
import constants
from data_loader import DataLoader, SplitDataLoader
sys.path.insert(0, '../src/modeling')
from network import Network

# Suppress TensorFlow messages
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '1' # 1 for INFO, 2 for INFO & WARNINGS, 3_
↳for INFO & WARNINGS & ERRORS
```

```
[2]: # visualization.py module cannot be imported because it uses non-GUI matplotlib
      ↪ backend.
      # Non-GUI backends does NOT allow to show figures, so new functions are needed
      ↪ for this notebook.

def plot_history(history, metric, output_path):

    """
    Plots the training history (train_loss vs val_loss, train_acc vs val_acc).

    Parameters:
        - history (dict):
            Train history.
        - metric (str):
            Metric to plot.
        - output_path (str):
            Absolute path to the .SVG file containing the plot.
    """

    f = plt.figure(figsize=(10,10))

    train_label = f'train_{metric}'
    val_label = f'val_{metric}'
    title = f'Train and Val {metric.title()}' # .title() upper-cases only the
    ↪ first letter

    plt.plot(history[metric], label=train_label)
    plt.plot(history[val_label], label=val_label)
    if metric == 'accuracy':
        plt.axhline(y=1/256, color='r', linewidth=3, linestyle='--',
    ↪ label='Random-Guesser Accuracy')
    plt.title(title)
    plt.ylabel(metric.title())
    plt.xlabel('Epochs')
    plt.legend()
    plt.grid()

    f.savefig(
        output_path,
        bbox_inches='tight',
        dpi=600
    )

    plt.show()

    plt.close(f)
```

```

def plot_conf_matrix(conf_matrix, output_path):

    cmap = plt.cm.Blues

    f = plt.figure(figsize=(10,8))

    plt.imshow(conf_matrix, cmap=cmap)
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')

    plt.colorbar()

    f.savefig(
        output_path,
        bbox_inches='tight',
        dpi=600
    )

    plt.show()

    plt.close(f)

def plot_ge(ge, output_path):

    """
    Plots the provided GE vector.

    Parameters:
        - ge (np.array):
            GE vector to plot.
        - output_path (str):
            Absolute path to the PNG file containing the plot.
    """

    # Plot GE
    f, ax = plt.subplots(figsize=(15,8))

    ax.plot(ge, marker='o', color='b')

    ax.set_title(f'Train Devices: D1, D2 | Test Device: D3')
    ax.set_xticks(range(len(ge)), labels=range(1, len(ge)+1))
    ax.set_xlabel('Number of traces')
    ax.set_ylabel('GE')
    ax.grid()

```

```

    f.savefig(
        output_path,
        bbox_inches='tight',
        dpi=600
    )

    plt.show()

    plt.close(f)

def plot_probs(prob_values, true_label, title, output_path):

    f = plt.figure(figsize=(15,8))

    for i in range(256):
        plt.plot([i, i], [0, prob_values[i]], color='b')
        plt.plot([i], [prob_values[i]], marker='o', color='b')

    plt.axvline(x=true_label, color='r', linestyle='--', label='True Label')

    plt.title(title)
    plt.xlabel('Labels')
    plt.ylabel('Prediction Probability')
    plt.legend(loc='upper right')
    plt.grid()

    plt.savefig(
        output_path,
        bbox_inches='tight',
        dpi=600
    )

    plt.show()

    plt.close(f)

```

```

[3]: RES_ROOT = f'{constants.RESULTS_PATH}/DKTA/SBOX_OUT/byte5/2d'
    HP_PATH = RES_ROOT + '/hp.json'

    SINGLE_ATTACK_FOLDER = f'{constants.RESULTS_PATH}/SingleAttack_Byte5'
    MODEL_PATH = SINGLE_ATTACK_FOLDER + f'/model_b5.h5'

    # Every plot has a .CSV file with data and a .svg file with the plot
    LOSS_HIST_FILE = SINGLE_ATTACK_FOLDER + '/loss_hist.csv'
    LOSS_HIST_PLOT = SINGLE_ATTACK_FOLDER + '/loss_hist.svg'

```

```

ACC_HIST_FILE = SINGLE_ATTACK_FOLDER + '/acc_hist.csv'
ACC_HIST_PLOT = SINGLE_ATTACK_FOLDER + '/acc_hist.svg'

CONF_MATRIX_FILE = SINGLE_ATTACK_FOLDER + '/conf_matrix.csv'
CONF_MATRIX_PLOT = SINGLE_ATTACK_FOLDER + '/conf_matrix.svg'

PROBS_MC_FILE = SINGLE_ATTACK_FOLDER + '/probs_mc.csv'
PROBS_MC_PLOT = SINGLE_ATTACK_FOLDER + '/probs_mc.svg'
PROBS_B_FILE = SINGLE_ATTACK_FOLDER + '/probs_b.csv'
PROBS_B_PLOT = SINGLE_ATTACK_FOLDER + '/probs_b.svg'
PROBS_W_FILE = SINGLE_ATTACK_FOLDER + '/probs_w.csv'
PROBS_W_PLOT = SINGLE_ATTACK_FOLDER + '/probs_w.svg'

GE_FILE = SINGLE_ATTACK_FOLDER + '/ge.csv'
GE_PLOT = SINGLE_ATTACK_FOLDER + '/ge.svg'

TRAIN_DEVS = ['D1', 'D2']
TEST_FILES = [f'{constants.PC_TRACES_PATH}/{d}-{k}_500MHz + Resampled.trs'] #_
↳list is needed in DataLoader

```

1.1 Training

```

[4]: # Get train data
train_files = [f'{constants.PC_TRACES_PATH}/{d}-{k}_500MHz + Resampled.trs'
               for k in list(constants.KEYS)[1:]
               for d in TRAIN_DEVS]

train_dl = SplitDataLoader(
    train_files,
    tot_traces=50000,
    train_size=0.9,
    target='SBOX_OUT',
    byte_idx=5
)
train_data, val_data = train_dl.load()
x_train, y_train, _, _ = train_data
x_val, y_val, _, _ = val_data

# Scale data to 0-mean and 1-variance
scaler = StandardScaler()
scaler.fit(x_train)
x_train = scaler.transform(x_train)
x_val = scaler.transform(x_val)

# Get hyperparameters

```

```

with open(HP_PATH, 'r') as jfile:
    hp = json.load(jfile)

# Train and save the model
net = Network('MLP', hp)
net.build_model()
net.add_checkpoint_callback(MODEL_PATH)

history = net.model.fit(
    x_train,
    y_train,
    validation_data=(x_val, y_val),
    epochs=100,
    batch_size=net.hp['batch_size'],
    callbacks=net.callbacks,
    verbose=0
).history

net.model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
Input (Dense)	(None, 1183)	1400672
InputBatchNorm (BatchNormalization)	(None, 1183)	4732
HiddenDropout0 (Dropout)	(None, 1183)	0
HiddenDense0 (Dense)	(None, 500)	592000
HiddenBatchNorm0 (BatchNormalization)	(None, 500)	2000
HiddenDropout1 (Dropout)	(None, 500)	0
HiddenDense1 (Dense)	(None, 500)	250500
HiddenBatchNorm1 (BatchNormalization)	(None, 500)	2000
HiddenDropout2 (Dropout)	(None, 500)	0
HiddenDense2 (Dense)	(None, 500)	250500

HiddenBatchNorm2 (BatchNorm alization)	(None, 500)	2000
HiddenDropout3 (Dropout)	(None, 500)	0
HiddenDense3 (Dense)	(None, 500)	250500
HiddenBatchNorm3 (BatchNorm alization)	(None, 500)	2000
OutputDropout (Dropout)	(None, 500)	0
Output (Dense)	(None, 256)	128256
OutputBatchNorm (BatchNorma alization)	(None, 256)	1024
Softmax (Activation)	(None, 256)	0

```
=====
Total params: 2,886,184
Trainable params: 2,879,306
Non-trainable params: 6,878
-----
```

The summary shows the **strucure of the model** and the **number of parameters per layer**.

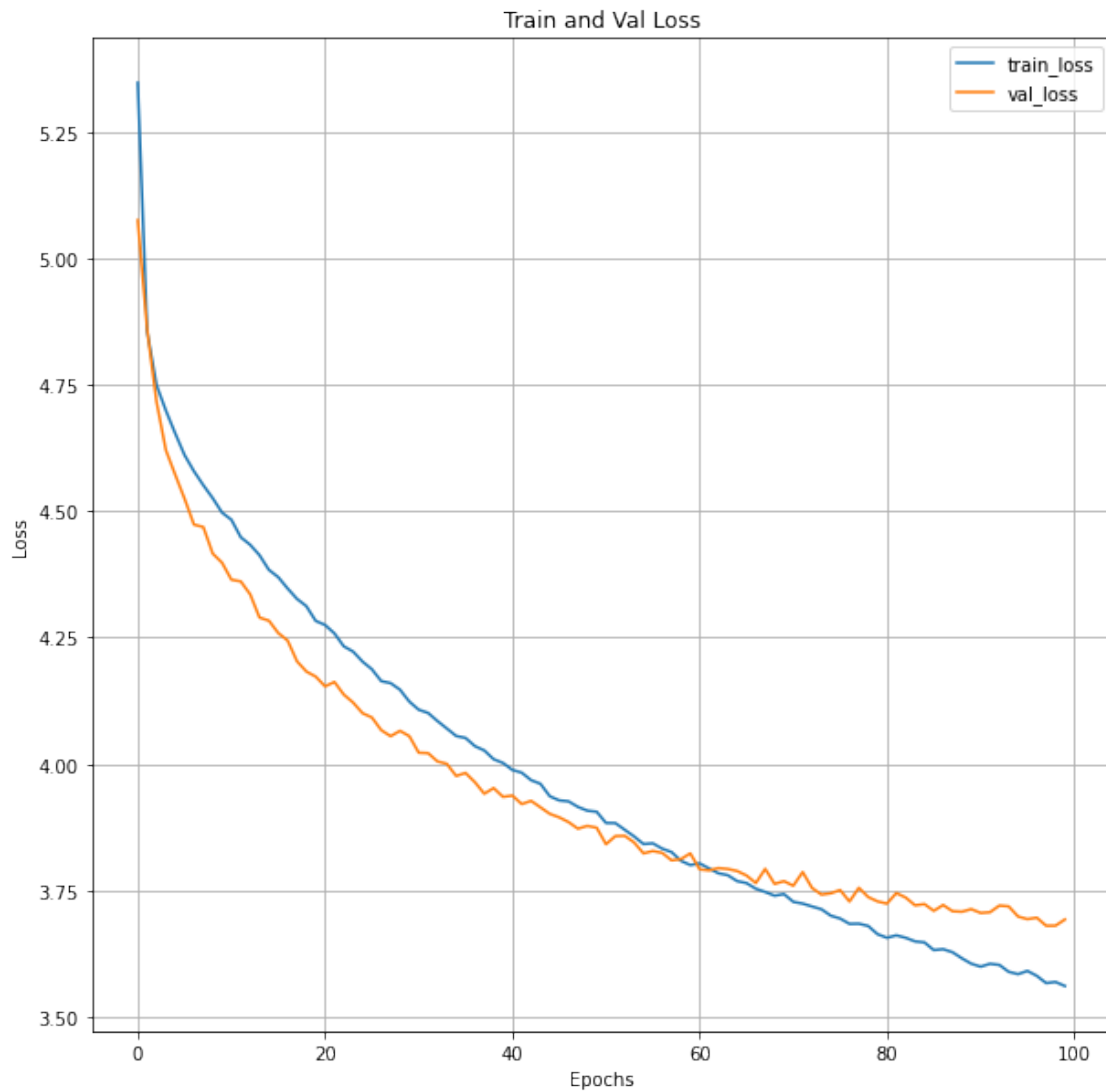
1.1.1 Training History

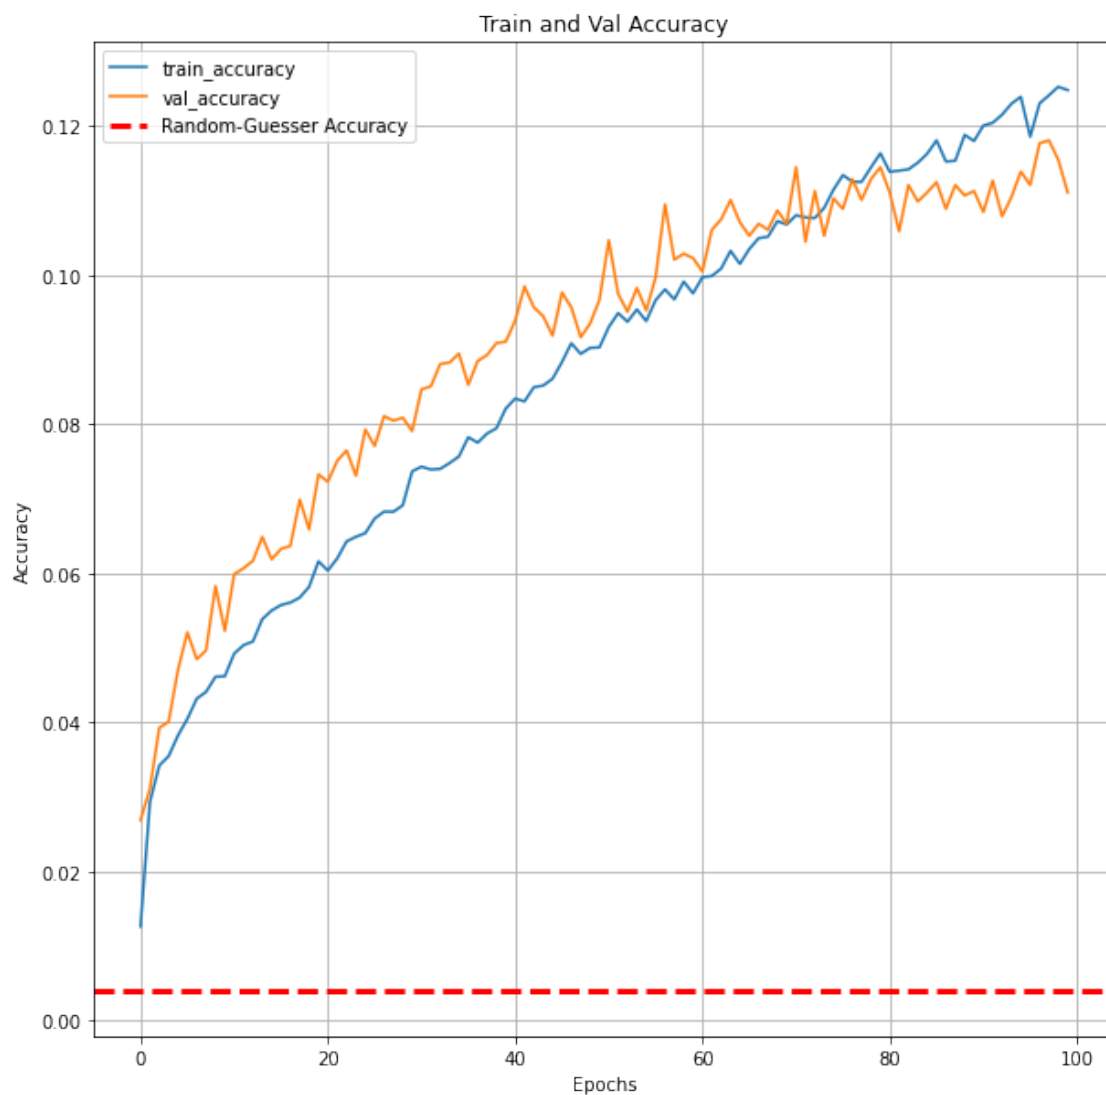
```
[5]: actual_epochs = len(history['loss']) # Early Stopping can make the actual
                                         # number of epochs different from the
                                         # original one
# Loss
loss_data = np.vstack(
    (
        np.arange(actual_epochs)+1, # X-axis values
        history['loss'], # Y-axis values for 'loss'
        history['val_loss'] # Y-axis values for 'val_loss'
    )
).T
helpers.save_csv(
    data=loss_data,
    columns=['Epochs', 'Loss', 'Val_Loss'],
    output_path=LOSS_HIST_FILE
)
plot_history(history, 'loss', LOSS_HIST_PLOT)
```

```

# Accuracy
acc_data = np.vstack(
    (
        np.arange(actual_epochs)+1, # X-axis values
        history['accuracy'], # Y-axis values for 'loss'
        history['val_accuracy'] # Y-axis values for 'val_loss'
    )
).T
helpers.save_csv(
    data=acc_data,
    columns=['Epochs', 'Acc', 'Val_Acc'],
    output_path=ACC_HIST_FILE
)
plot_history(history, 'accuracy', ACC_HIST_PLOT)

```





Validation-metrics are better than train-metrics for a while. However, at the end of the training process, the validation-metrics are slightly worse than the train-metrics.

This is a strange behavior, most probably caused by the following reasons: * **Regularization:** it is applied during training, while it is removed during the validation step, leading to a penalized model during training and a potentially better one for validation (in this case Dropout is applied with a 40% chance); * **Metrics computation:** train-metrics are computed *during* each epoch, while validation-metrics are measured *after* each epoch, resulting in a slight misalignment.

However, **validation-metrics do not diverge from train-metrics**, showing that **the model is not overfitting** the data.

The model shows also a very good Validation Accuracy for the problem: **the model outperforms a Random-Guesser**, whose accuracy is $1/256 \sim 0.0039$.

1.2 Attack

```
[6]: attack_dl = DataLoader(
    TEST_FILES,
    tot_traces=50000,
    target='SBOX_OUT',
    byte_idx=5
)
x_attack, y_attack, pbs_attack, tkb_attack = attack_dl.load()

# Scale data to 0-mean and 1-variance
x_attack = scaler.transform(x_attack)

attack_model = load_model(MODEL_PATH)
```

1.2.1 Attack Loss & Attack Accuracy

```
[7]: attack_loss, attack_acc = attack_model.evaluate(x_attack, y_attack, verbose=0)

print(f'Attack Loss:      {attack_loss:.2f}')
print(f'Attack Accuracy: {(attack_acc*100):.2f}%')
```

```
Attack Loss:      4.67
Attack Accuracy: 5.57%
```

The obtained results are not the best for a classical classification problem, but they are **enough for this specific scenario**.

Indeed, in Deep Learning based SCAs there is the additional step of “combining the single predictions”, because all the input data share the true objective of the attack, the key.

However, it is interesting to see the values of the metrics because both hyperparameter tuning and training callbacks are based on them, specifically on validation loss.

1.2.2 Confusion Matrix

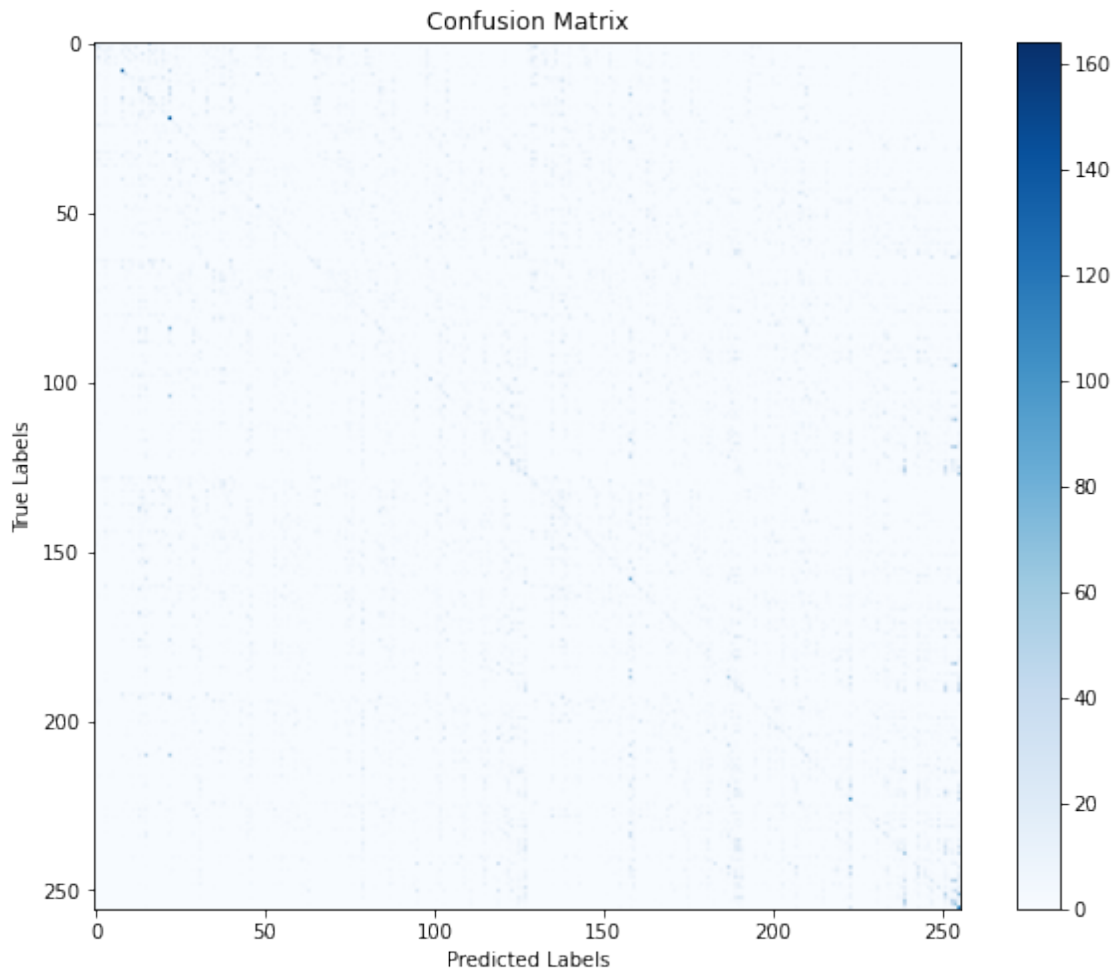
```
[8]: preds = attack_model.predict(x_attack)

# From one-hot-encoding to true labels in [0;255]
y_true = [el.tolist().index(1) for el in y_attack]
y_pred = [el.tolist().index(max(el)) for el in preds]

cm_data = confusion_matrix(y_true, y_pred)

helpers.save_csv(
    data=cm_data,
    columns=[f'Col_{i+1}' for i in range(cm_data.shape[1])],
    output_path=CONF_MATRIX_FILE
)
```

```
plot_conf_matrix(cm_data, CONF_MATRIX_PLOT)
```



The ideal confusion matrix has values only on the main diagonal (each predicted label matches the corresponding true label).

In this case there are misclassifications, but **the main diagonal is still visible**, meaning that the model is not a random guesser.

The **misclassifications are consequence of the low attack accuracy** (~11%).
However, **this result is good enough to allow the correct key-byte recovery**.

1.2.3 Guessing Entropy

```
[9]: ge = results.ge(  
      model=attack_model,  
      x_test=x_attack,  
      pltxt_bytes=pbs_attack,
```

```

    true_key_byte=tkb_attack,
    n_exp=100,
    target='SBOX_OUT'
)

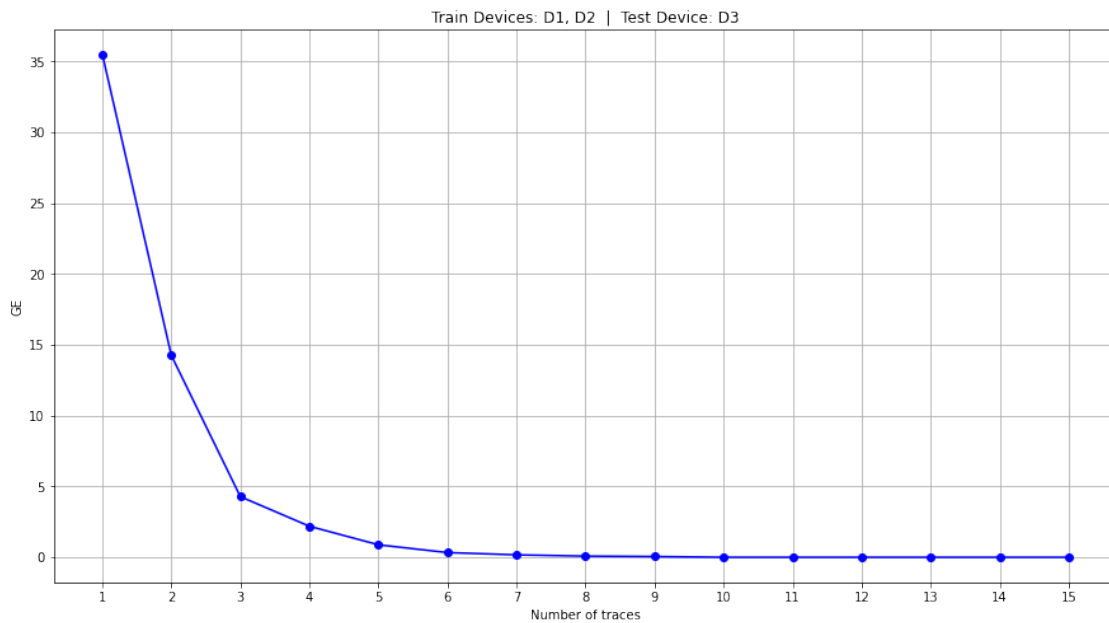
plottable_ge = ge[:15]

csv_ge_data = np.vstack(
    (
        np.arange(plottable_ge.shape[0])+1,
        plottable_ge
    )
).T

helpers.save_csv(
    data=csv_ge_data,
    columns=['NTraces', 'GE'],
    output_path=GE_FILE
)

plot_ge(plottable_ge, GE_PLOT)

```



The overall performance is very good: **GE reaches 0 in less than 10 traces**, allowing the recovery of the 6th byte of the key considering only 6 attack traces (on average).

This result shows how in Deep Learning based SCAs accuracy should not be the reference metric: **a model with high accuracy will surely perform well, but a model with low accuracy**

cannot be discarded a priori.

1.2.4 Byte Recovery

```
[30]: import random

ATTACKED_KEY = np.array([int(kb, 16) for kb in constants.KEYS['K0']])
BYTE_6 = ATTACKED_KEY[5]
N_EXP = 10
N_TR_PER_EXP = 10

print(f'Correct 6th Byte: {BYTE_6}')
print()
print(f'Performing {N_EXP} Experiments considering {N_TR_PER_EXP} attack-traces.
↪...')
print()

attack_set = list(zip(x_attack, pbs_attack))
attack_subset = random.sample(attack_set, N_EXP*N_TR_PER_EXP) # Get a random
↪set of attack traces (with relative plaintexts)
selected_x, selected_pbs = zip(*attack_subset)
selected_x = np.vstack(selected_x)
selected_pbs = np.vstack(selected_pbs)

for i in range(N_EXP):
    print(f'Exp {i+1}')
    start = i * N_TR_PER_EXP
    stop = start + N_TR_PER_EXP

    x_batch = selected_x[start:stop]
    pbs_batch = selected_pbs[start:stop]

    curr_preds = attack_model.predict(x_batch)

    predicted_byte_inc_traces = results.retrieve_key_byte(
        preds=curr_preds,
        pltxt_bytes=pbs_batch,
        target='SBOX_OUT'
    ).tolist()

    try:
        rank = predicted_byte_inc_traces.index(BYTE_6)
        print(f'Correct byte found in {rank+1} attack traces!')
    except:
        print('Error: Correct byte NOT found')
```

Correct 6th Byte: 253

Performing 10 Experiments considering 10 attack-traces...

Exp 1
Correct byte found in 4 attack traces!
Exp 2
Correct byte found in 4 attack traces!
Exp 3
Correct byte found in 8 attack traces!
Exp 4
Correct byte found in 1 attack traces!
Exp 5
Correct byte found in 4 attack traces!
Exp 6
Correct byte found in 4 attack traces!
Exp 7
Correct byte found in 8 attack traces!
Exp 8
Correct byte found in 2 attack traces!
Exp 9
Correct byte found in 2 attack traces!
Exp 10
Correct byte found in 3 attack traces!

The attacked byte is **correctly recovered in all experiments**.

1.2.5 Low Accuracy Considerations

```
[ ]: # MC = (6, 'Most Common Case')
      # B = (42, 'Best Case')
      # W = (41, 'Worst Case')

      # CASES = [MC, B, W]
      # OUTPUT_FILES = [PROBS_MC_FILE, PROBS_B_FILE, PROBS_W_FILE]
      # OUTPUT_PLOTS = [PROBS_MC_PLOT, PROBS_B_PLOT, PROBS_W_PLOT]

      # for i, c in enumerate(CASES):

      #     tr, title = c

      #     print(title)

      #     csv_probs_data = np.vstack(
      #         (
      #             np.arange(256),
      #             preds[tr],
      #             y_attack[tr]
      #         )
      #     ).T
```

```
#     helpers.save_csv(  
#         data=csv_probs_data,  
#         columns=['Label', 'Predicted_Prob', 'Actual_Prob'],  
#         output_path=OUTPUT_FILES[i]  
#     )  
  
#     true_label = y_attack[tr].tolist().index(1)  
#     plot_probs(preds[tr], true_label, title, OUTPUT_PLOTS[i])
```